

```
In [5]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

pd.set_option('display.max_columns', None)
plt.rcParams['figure.figsize'] = (12,6)
```

```
In [4]: data = pd.read_csv("E:\data set\Phishing_Legitimate_full.csv")
```

```
In [3]: float_cols = data.select_dtypes('float64').columns
for c in float_cols:
    data[c] = data[c].astype('float32')

int_cols = data.select_dtypes('int64').columns
for c in int_cols:
    data[c] = data[c].astype('int32')

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 50 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   id                                         10000 non-null  int32
1   NumDots                                   10000 non-null  int32
2   SubdomainLevel                           10000 non-null  int32
3   PathLevel                                10000 non-null  int32
4   UrlLength                                10000 non-null  int32
5   NumDash                                   10000 non-null  int32
6   NumDashInHostname                       10000 non-null  int32
7   AtSymbol                                 10000 non-null  int32
8   TildeSymbol                             10000 non-null  int32
9   NumUnderscore                           10000 non-null  int32
10  NumPercent                               10000 non-null  int32
11  NumQueryComponents                       10000 non-null  int32
12  NumAmpersand                             10000 non-null  int32
13  NumHash                                  10000 non-null  int32
14  NumNumericChars                         10000 non-null  int32
15  NoHttps                                  10000 non-null  int32
16  RandomString                             10000 non-null  int32
17  IpAddress                                10000 non-null  int32
18  DomainInSubdomains                      10000 non-null  int32
19  DomainInPaths                           10000 non-null  int32
20  HttpsInHostname                         10000 non-null  int32
21  HostnameLength                          10000 non-null  int32
22  PathLength                              10000 non-null  int32
23  QueryLength                             10000 non-null  int32
24  DoubleSlashInPath                      10000 non-null  int32
25  NumSensitiveWords                       10000 non-null  int32
26  EmbeddedBrandName                      10000 non-null  int32
27  PctExtHyperlinks                       10000 non-null  float32
28  PctExtResourceUrls                     10000 non-null  float32
29  ExtFavicon                              10000 non-null  int32
30  InsecureForms                          10000 non-null  int32
31  RelativeFormAction                     10000 non-null  int32
32  ExtFormAction                          10000 non-null  int32
33  AbnormalFormAction                     10000 non-null  int32
34  PctNullSelfRedirectHyperlinks          10000 non-null  float32
35  FrequentDomainNameMismatch             10000 non-null  int32
36  FakeLinkInStatusBar                    10000 non-null  int32
37  RightClickDisabled                     10000 non-null  int32
38  PopUpWindow                            10000 non-null  int32
39  SubmitInfoToEmail                      10000 non-null  int32
40  IFrameOrFrame                          10000 non-null  int32
41  MissingTitle                           10000 non-null  int32
42  ImagesOnlyInForm                       10000 non-null  int32
43  SubdomainLevelRT                       10000 non-null  int32
44  UrlLengthRT                             10000 non-null  int32
45  PctExtResourceUrlsRT                   10000 non-null  int32
46  AbnormalExtFormActionR                  10000 non-null  int32
47  ExtMetaScriptLinkRT                    10000 non-null  int32
48  PctExtNullSelfRedirectHyperlinksRT      10000 non-null  int32
49  CLASS_LABEL                             10000 non-null  int32
dtypes: float32(3), int32(47)
memory usage: 1.9 MB
```

```
In [7]: data.rename(columns={'CLASS_LABEL': 'labels'}, inplace=True)
```

```
In [8]: data.sample(5)
```


Out[8]:											
	id	NumDots	SubdomainLevel	PathLevel	UrlLength	NumDash	NumDashInHostname	AtSymbol	TildeSymbol	NumUnderscore	Num
6824	6825	3	1	1	103	10	3	0	0	0	
9438	9439	2	1	3	45	0	0	0	0	0	
2826	2827	3	0	5	62	1	0	0	0	0	
2580	2581	4	1	1	42	1	1	0	0	0	
5354	5355	2	1	5	134	6	0	0	0	2	

```
In [9]: data.describe()
```

Out[9]:		id	NumDots	SubdomainLevel	PathLevel	UrlLength	NumDash	NumDashInHostname	AtSymbol	TildeSyn
	count	10000.00000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000
	mean	5000.50000	2.445100	0.586800	3.300300	70.264100	1.818000	0.138900	0.000300	0.013000
	std	2886.89568	1.346836	0.751214	1.863241	33.369877	3.106258	0.545744	0.017319	0.113000
	min	1.00000	1.000000	0.000000	0.000000	12.000000	0.000000	0.000000	0.000000	0.000000
	25%	2500.75000	2.000000	0.000000	2.000000	48.000000	0.000000	0.000000	0.000000	0.000000
	50%	5000.50000	2.000000	1.000000	3.000000	62.000000	0.000000	0.000000	0.000000	0.000000
	75%	7500.25000	3.000000	1.000000	4.000000	84.000000	2.000000	0.000000	0.000000	0.000000
	max	10000.00000	21.000000	14.000000	18.000000	253.000000	55.000000	9.000000	1.000000	1.000000

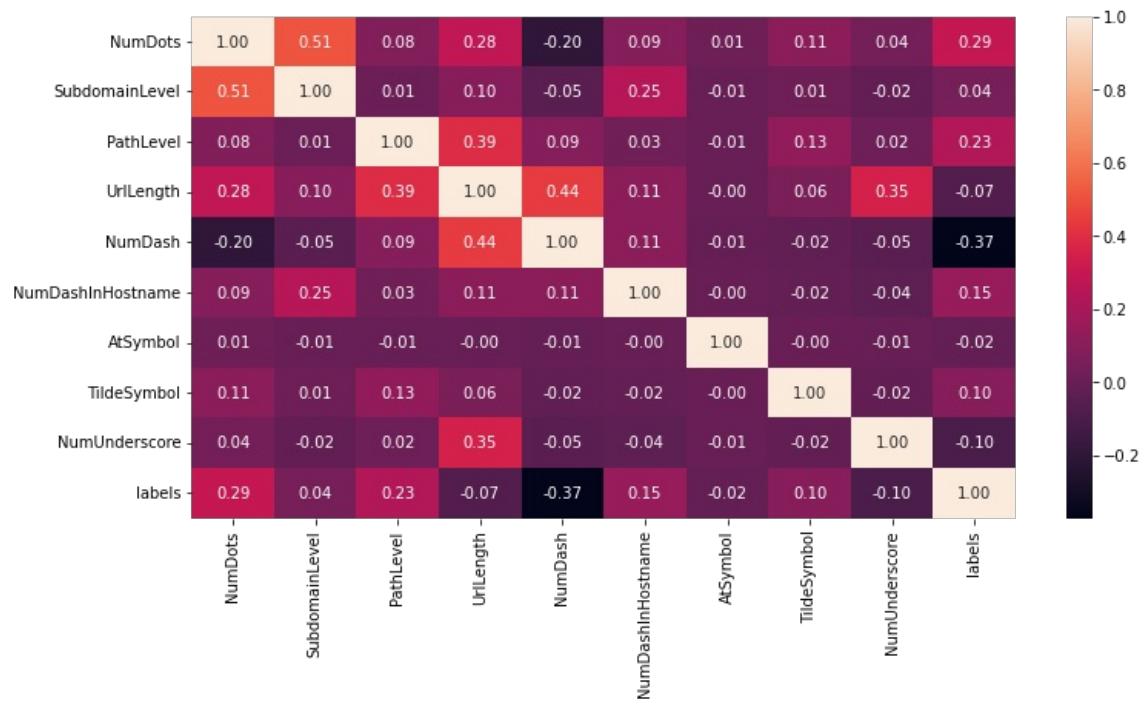
```
In [10]: data['labels'].value_counts().plot(kind='bar')
```

Out[10]: <AxesSubplot:>

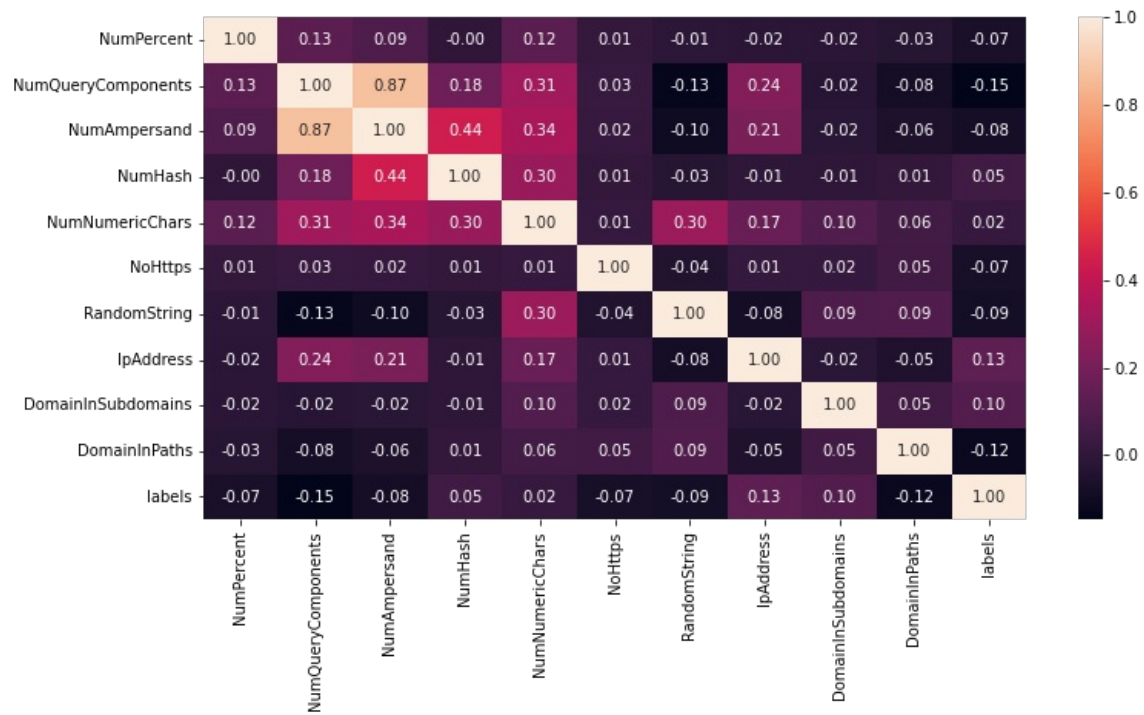


```
In [11]: def corr_heatmap(data, idx_s, idx_e):
          y = data['labels']
          temp = data.iloc[:, idx_s:idx_e]
          if 'id' in temp.columns:
              del temp['id']
          temp['labels'] = y
          sns.heatmap(temp.corr(), annot=True, fmt='.2f')
          plt.show()
```

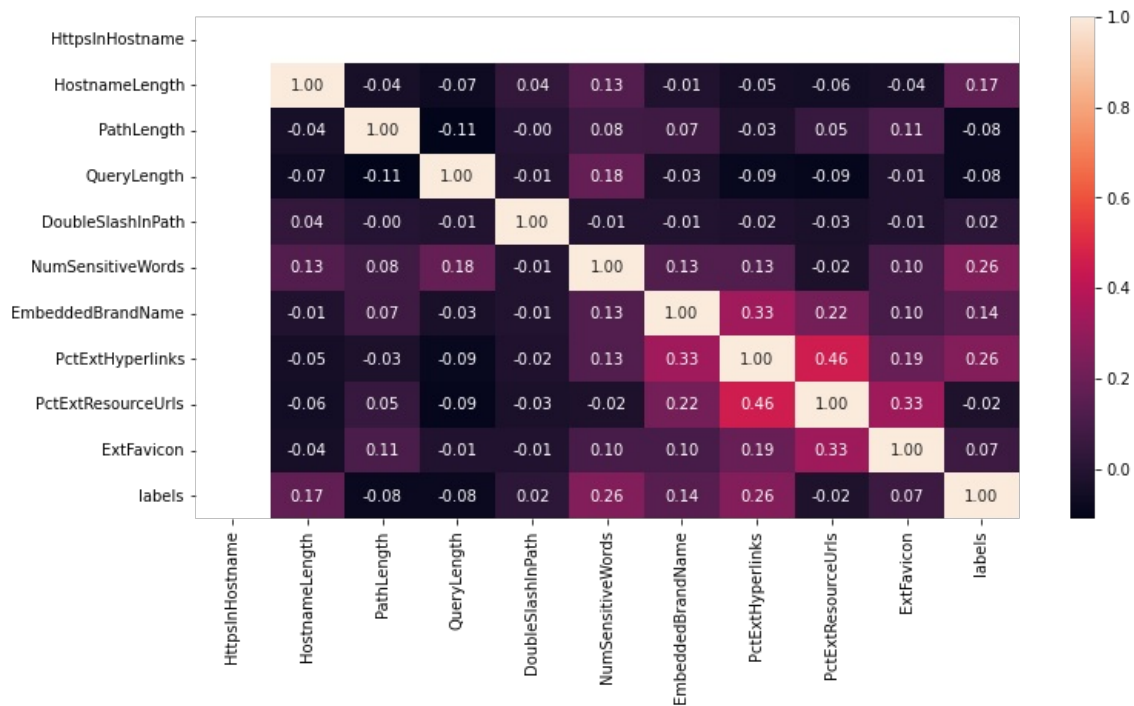
```
In [12]: corr_heatmap(data, 0, 10)
```



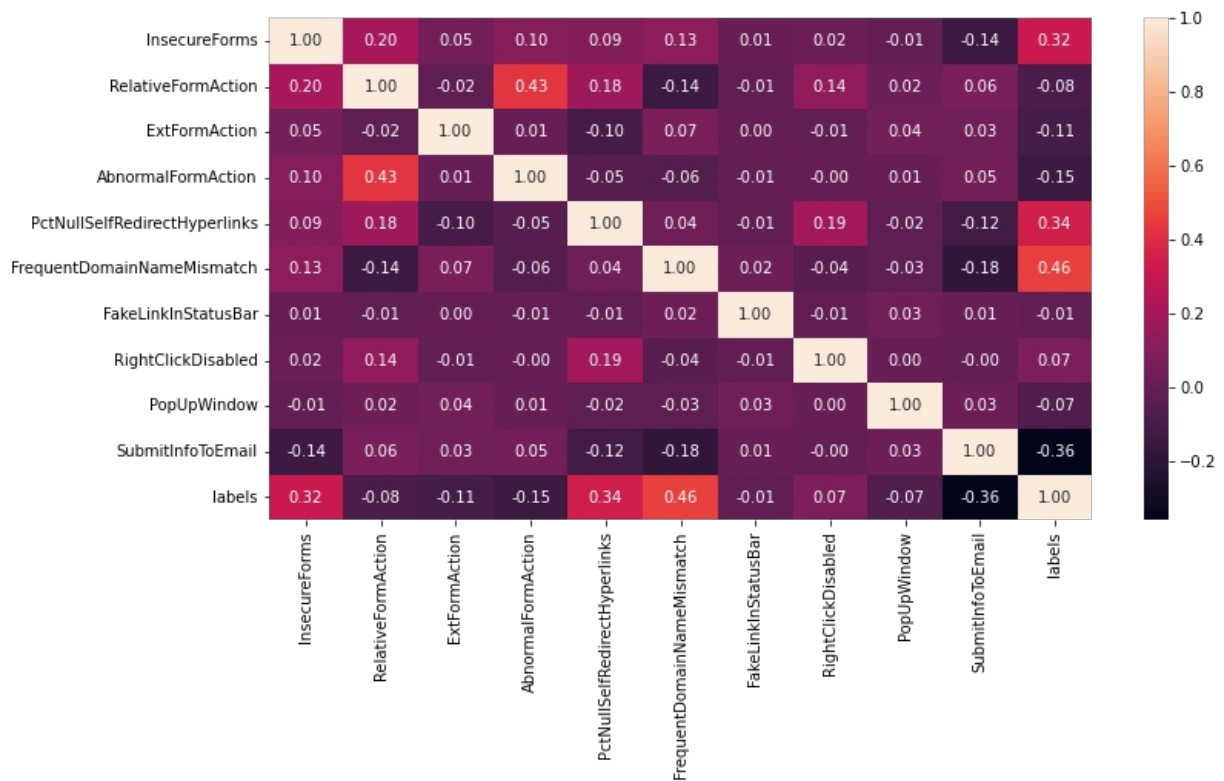
```
In [13]: corr_heatmap(data, 10, 20)
```



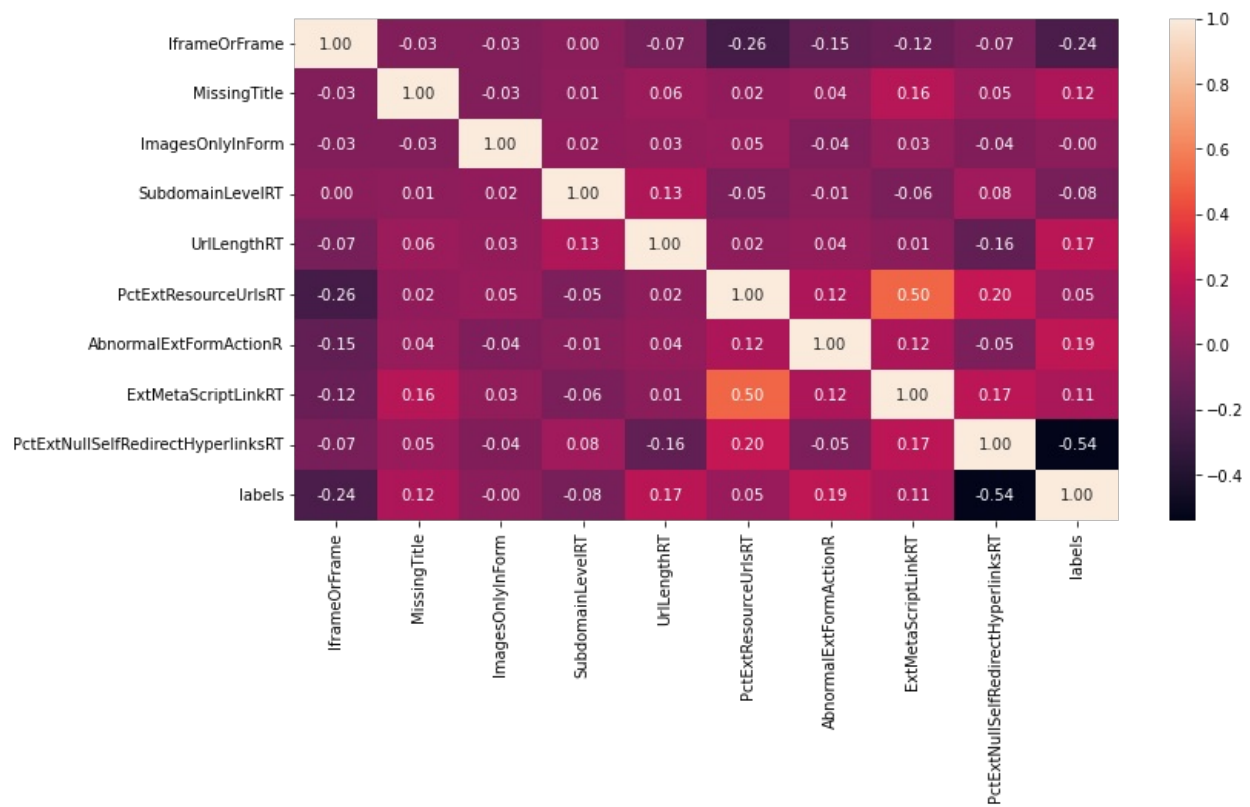
```
In [14]: corr_heatmap(data, 20, 30)
```



```
In [15]: corr_heatmap(data, 30, 40)
```



```
In [16]: corr_heatmap(data, 40, 50)
```



```
In [17]: from sklearn.feature_selection import mutual_info_classif
```

```
In [18]: X = data.drop(['id', 'labels'], axis=1)
y = data['labels']
```


```
In [19]: discrete_features = X.dtypes == int
```

```
In [20]: mi_scores = mutual_info_classif(X, y, discrete_features=discrete_features)
mi_scores = pd.Series(mi_scores, name="MI Scores", index=X.columns)
mi_scores = mi_scores.sort_values(ascending=False)
mi_scores
```

```
Out[20]: PctExtHyperlinks      0.470615
PctExtResourceUrls      0.287998
PctNullSelfRedirectHyperlinks 0.233576
PctExtNullSelfRedirectHyperlinksRT 0.212144
FrequentDomainNameMismatch 0.131566
NumNumericChars        0.129133
ExtMetaScriptLinkRT     0.121971
NumDash                 0.112274
SubmitInfoToEmail       0.080130
NumDots                 0.071204
InsecureForms           0.062497
PathLevel               0.054670
QueryLength             0.054056
PathLength              0.047413
NumSensitiveWords       0.038090
UrlLength               0.032149
HostnameLength          0.030030
NumQueryComponents      0.027929
PctExtResourceUrlsRT    0.021614
NumAmpersand            0.020247
AbnormalExtFormActionR  0.015332
UrlLengthRT             0.015089
NumDashInHostname       0.014805
IpAddress               0.013922
MissingTitle            0.013292
AbnormalFormAction      0.012685
IframeOrFrame           0.012480
ExtFormAction           0.011569
NumHash                 0.011511
EmbeddedBrandName       0.011107
NumPercent              0.010437
NumUnderscore           0.007182
PopUpWindow             0.006713
TildeSymbol             0.005063
DomainInPaths           0.004764
RelativeFormAction      0.004580
SubdomainLevelRT        0.004332
FakeLinkInStatusBar     0.003984
SubdomainLevel          0.002838
RightClickDisabled      0.002098
RandomString            0.002029
DomainInSubdomains      0.001879
ExtFavicon              0.001787
HttpsInHostname         0.000000
ImagesOnlyInForm        0.000000
NoHttps                 0.000000
AtSymbol                0.000000
DoubleSlashInPath       0.000000
Name: MI Scores, dtype: float64
```

```
In [21]: def plot_mi_scores(scores):
scores = scores.sort_values(ascending=True)
width = np.arange(len(scores))
ticks = list(scores.index)
plt.barh(width, scores)
plt.yticks(width, ticks)
plt.title("MI Scores")

plt.figure(dpi=100, figsize=(12,12))
plot_mi_scores(mi_scores)
```



```
In [25]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```
In [26]: def train_logistic(data, top_n):
top_n_features = mi_scores.sort_values(ascending=False).head(top_n).index.tolist()
X = data[top_n_features]
y = data['labels']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=True)

lr = LogisticRegression(max_iter=10000)
lr.fit(X_train, y_train)

y_pred = lr.predict(X_test)

precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)

return precision, recall, f1, accuracy
```

```
In [27]: arr = []
for i in range(20,51,1):
precision, recall, f1, accuracy = train_logistic(data, i)
print("Performance for Logistic Model with Top {} features is precision : {}, recall : {}, f1 score : {}, a
arr.append([i, precision, recall, f1, accuracy])
```


Performance for Logistic Model with Top 20 features is precision : 0.8943248532289628, recall : 0.9298067141403
 866, f1 score : 0.9117206982543641, accuracy : 0.9115
 Performance for Logistic Model with Top 21 features is precision : 0.906496062992126, recall : 0.94364754098360
 66, f1 score : 0.924698795180723, accuracy : 0.925
 Performance for Logistic Model with Top 22 features is precision : 0.9306243805748265, recall : 0.939, f1 score :
 : 0.9347934295669488, accuracy : 0.9345
 Performance for Logistic Model with Top 23 features is precision : 0.9166666666666666, recall : 0.9354508196721
 312, f1 score : 0.9259634888438133, accuracy : 0.927
 Performance for Logistic Model with Top 24 features is precision : 0.9279835390946503, recall : 0.9376299376299
 376, f1 score : 0.9327817993795242, accuracy : 0.935
 Performance for Logistic Model with Top 25 features is precision : 0.9195289499509323, recall : 0.9398194583751
 254, f1 score : 0.9295634920634921, accuracy : 0.929
 Performance for Logistic Model with Top 26 features is precision : 0.9164208456243854, recall : 0.9481180061037
 64, f1 score : 0.932, accuracy : 0.932
 Performance for Logistic Model with Top 27 features is precision : 0.9349112426035503, recall : 0.9423459244532
 804, f1 score : 0.9386138613861386, accuracy : 0.938
 Performance for Logistic Model with Top 28 features is precision : 0.9434914228052472, recall : 0.9285004965243
 296, f1 score : 0.9359359359359359, accuracy : 0.936
 Performance for Logistic Model with Top 29 features is precision : 0.9255533199195171, recall : 0.9255533199195
 171, f1 score : 0.9255533199195171, accuracy : 0.926
 Performance for Logistic Model with Top 30 features is precision : 0.9256360078277887, recall : 0.9422310756972
 112, f1 score : 0.9338598223099704, accuracy : 0.933
 Performance for Logistic Model with Top 31 features is precision : 0.9165867689357622, recall : 0.9446640316205
 533, f1 score : 0.9304136253041362, accuracy : 0.9285
 Performance for Logistic Model with Top 32 features is precision : 0.9058252427184466, recall : 0.9339339339339
 34, f1 score : 0.9196648595367176, accuracy : 0.9185
 Performance for Logistic Model with Top 33 features is precision : 0.9258188824662813, recall : 0.9533730158730
 159, f1 score : 0.9393939393939394, accuracy : 0.938
 Performance for Logistic Model with Top 34 features is precision : 0.9299691040164778, recall : 0.9357512953367
 876, f1 score : 0.9328512396694214, accuracy : 0.935
 Performance for Logistic Model with Top 35 features is precision : 0.927710843373494, recall : 0.93807106598984
 77, f1 score : 0.9328621908127208, accuracy : 0.9335
 Performance for Logistic Model with Top 36 features is precision : 0.9295366795366795, recall : 0.9563058589870
 904, f1 score : 0.9427312775330398, accuracy : 0.9415
 Performance for Logistic Model with Top 37 features is precision : 0.9315332690453231, recall : 0.9498525073746
 312, f1 score : 0.9406037000973709, accuracy : 0.939
 Performance for Logistic Model with Top 38 features is precision : 0.916023166023166, recall : 0.94146825396825
 4, f1 score : 0.9285714285714286, accuracy : 0.927
 Performance for Logistic Model with Top 39 features is precision : 0.9408866995073891, recall : 0.9408866995073
 891, f1 score : 0.9408866995073891, accuracy : 0.94
 Performance for Logistic Model with Top 40 features is precision : 0.9409409409409409, recall : 0.9390609390609
 39, f1 score : 0.94, accuracy : 0.94
 Performance for Logistic Model with Top 41 features is precision : 0.9244357212953876, recall : 0.9448345035105
 316, f1 score : 0.9345238095238094, accuracy : 0.934
 Performance for Logistic Model with Top 42 features is precision : 0.9257425742574258, recall : 0.9415911379657
 603, f1 score : 0.9335996005991014, accuracy : 0.9335
 Performance for Logistic Model with Top 43 features is precision : 0.9464469618949537, recall : 0.9484004127966
 976, f1 score : 0.9474226804123711, accuracy : 0.949
 Performance for Logistic Model with Top 44 features is precision : 0.9303921568627451, recall : 0.9396039603960
 396, f1 score : 0.9349753694581281, accuracy : 0.934
 Performance for Logistic Model with Top 45 features is precision : 0.936105476673428, recall : 0.94958847736625
 52, f1 score : 0.9427987742594486, accuracy : 0.944
 Performance for Logistic Model with Top 46 features is precision : 0.9460539460539461, recall : 0.9527162977867
 203, f1 score : 0.94937343358396, accuracy : 0.9495
 Performance for Logistic Model with Top 47 features is precision : 0.9277942631058358, recall : 0.9551934826883
 91, f1 score : 0.9412945308580031, accuracy : 0.9415
 Performance for Logistic Model with Top 48 features is precision : 0.9484848484848485, recall : 0.9361914257228
 315, f1 score : 0.9422980431510286, accuracy : 0.9425
 Performance for Logistic Model with Top 49 features is precision : 0.9375600384245918, recall : 0.9457364341085
 271, f1 score : 0.9416304872165943, accuracy : 0.9395
 Performance for Logistic Model with Top 50 features is precision : 0.937984496124031, recall : 0.94901960784313
 72, f1 score : 0.9434697855750486, accuracy : 0.942

```
In [28]: df = pd.DataFrame(arr, columns=['num_of_features', 'precision', 'recall', 'f1_score', 'accuracy'])
df
```




```
Out[28]:
```

	num_of_features	precision	recall	f1_score	accuracy
0	20	0.894325	0.929807	0.911721	0.9115
1	21	0.906496	0.943648	0.924699	0.9250
2	22	0.930624	0.939000	0.934793	0.9345
3	23	0.916667	0.935451	0.925963	0.9270
4	24	0.927984	0.937630	0.932782	0.9350
5	25	0.919529	0.939819	0.929563	0.9290
6	26	0.916421	0.948118	0.932000	0.9320
7	27	0.934911	0.942346	0.938614	0.9380
8	28	0.943491	0.928500	0.935936	0.9360
9	29	0.925553	0.925553	0.925553	0.9260
10	30	0.925636	0.942231	0.933860	0.9330
11	31	0.916587	0.944664	0.930414	0.9285
12	32	0.905825	0.933934	0.919665	0.9185
13	33	0.925819	0.953373	0.939394	0.9380
14	34	0.929969	0.935751	0.932851	0.9350
15	35	0.927711	0.938071	0.932862	0.9335
16	36	0.929537	0.956306	0.942731	0.9415
17	37	0.931533	0.949853	0.940604	0.9390
18	38	0.916023	0.941468	0.928571	0.9270
19	39	0.940887	0.940887	0.940887	0.9400
20	40	0.940941	0.939061	0.940000	0.9400
21	41	0.924436	0.944835	0.934524	0.9340
22	42	0.925743	0.941591	0.933600	0.9335
23	43	0.946447	0.948400	0.947423	0.9490
24	44	0.930392	0.939604	0.934975	0.9340
25	45	0.936105	0.949588	0.942799	0.9440
26	46	0.946054	0.952716	0.949373	0.9495
27	47	0.927794	0.955193	0.941295	0.9415
28	48	0.948485	0.936191	0.942298	0.9425
29	49	0.937560	0.945736	0.941630	0.9395
30	50	0.937984	0.949020	0.943470	0.9420

```
In [29]: sns.lineplot(x='num_of_features', y='precision', data=df, label='Precision Score')
sns.lineplot(x='num_of_features', y='recall', data=df, label='Recall Score')
sns.lineplot(x='num_of_features', y='f1_score', data=df, label='F1 Score')
sns.lineplot(x='num_of_features', y='accuracy', data=df, label='Acc Score')
```

```
Out[29]: <AxesSubplot:xlabel='num_of_features', ylabel='precision'>
```



```
In [30]: def train_rfc(data, top_n):
top_n_features = mi_scores.sort_values(ascending=False).head(top_n).index.tolist()
X = data[top_n_features]
y = data['labels']
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=True)

rfc = cuRfc(n_estimators=500,
            split_criterion=1,
            max_depth=32,
            max_leaves=-1,
            max_features=1.0,
            n_bins=128)

rfc.fit(X_train, y_train)

y_pred = rfc.predict(X_test, predict_model='CPU')

precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)

return precision, recall, f1, accuracy

```

```

In [33]: df = pd.DataFrame(arr, columns=['num_of_features', 'precision', 'recall', 'f1_score', 'accuracy'])
df.head()

```

```

Out[33]:
   num_of_features  precision  recall  f1_score  accuracy
0                20    0.937984    0.94902    0.94347    0.942
1                21    0.937984    0.94902    0.94347    0.942
2                22    0.937984    0.94902    0.94347    0.942
3                23    0.937984    0.94902    0.94347    0.942
4                24    0.937984    0.94902    0.94347    0.942

```

```

In [43]: from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
x_train=scaler.fit_transform(X_train)
x_test=scaler.transform(X_test)

```

```

In [45]: from sklearn.ensemble import RandomForestClassifier
classifier=RandomForestClassifier(random_state = 0)
classifier.fit(x_train, y_train)

```

```

Out[45]: RandomForestClassifier(random_state=0)

```

```

In [46]: y_pred1=classifier.predict(x_test)
y_pred1_train=classifier.predict(x_train)

```

```

In [54]: from sklearn.metrics import confusion_matrix,mean_absolute_error,mean_squared_error,recall_score,accuracy_score

```

```

In [48]: confusion_matrix(y_test,y_pred1)

```

```

Out[48]: array([[ 970,   14],
               [   14, 1002]], dtype=int64)

```

```

In [49]: print("Recall is", recall_score(y_test,y_pred1))

Recall is 0.9862204724409449

```

```

In [50]: print("Precision is", precision_score(y_test,y_pred1))

Precision is 0.9862204724409449

```

```

In [51]: print("Mean Absolute Error is",mean_absolute_error(y_test,y_pred1))

Mean Absolute Error is 0.014

```

```

In [52]: print("Mean Squared Error is",mean_squared_error(y_test,y_pred1))

Mean Squared Error is 0.014

```

```

In [55]: print(" Root Mean Squared Error is",mean_squared_error(y_test,y_pred1,squared=False))

Root Mean Squared Error is 0.11832159566199232

```

```

In [57]: print("Accuracy is",accuracy_score(y_test,y_pred1))

Accuracy is 0.986

```

```

In [59]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred1))

```

	precision	recall	f1-score	support
0	0.99	0.99	0.99	984
1	0.99	0.99	0.99	1016
accuracy			0.99	2000
macro avg	0.99	0.99	0.99	2000
weighted avg	0.99	0.99	0.99	2000

In []:

In []:

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js